

Introduction to Computer Systems, Spring 2023 (DaSE, ECNU)

Lab Assignment L4: Understanding Cache Memories

Assigned: Friday May 19, Due: Friday June 12 23:59 (UTC+8)

Harry Bovik (bovik@cs.cmu.edu) is the lead person for designing this lab.

Tong-yu Liu (tyliu@stu.ecnu.edu.cn) is the TA responsible for assigning and grading this lab.

1 Logistics

This is an individual project. You must complete and run this lab in your container of Shuishan.

2 Overview

This lab will help you understand the impact that cache memories can have on the performance of your C programs. You will write a small C program (about 200-300 lines) that simulates the behavior of a cache memory.

3 Initialization for the Assignment

This assignment should be submitted to your repository in Shuishan, so please follow the instructions below to initialize the experimental environment in your container first.

1. Download `cachelab-handout.tar` from the course webpage.
2. Copy `cachelab-handout.tar` to a protected directory in your container, where you plan to do your work.
3. Enter the directory and initialize the git repository for the assignment.

```
linux> cd <DIR>
linux> git init
```

4. Set the remote repository and pull the branch `homework04` from the remote.

```
linux> git remote add <REMOTE_NAME> <REMOTE_URL>
linux> git pull <REMOTE_NAME> homework04:master
```

At this point, the file `README.md` will appear in your directory. You should write your report in this file.

5. Decompress the tar file and (optionally) delete it.

```
linux> tar xvf cachelab-handout.tar
linux> rm -f cachelab-handout.tar
```

This will create a directory called `cachelab-handout` that contains a number of files. You will be modifying the file `csim.c`.

6. Stage the changes and make a commit.

```
linux> git add .
linux> git commit -m "initialize cachelab"
```

After completing the above steps, you can complete your assignment in the repository. Remember, you should often use `git commit` to save your progress.

In the directory `cachelab-handout`, you can compile all files by typing

```
linux> make clean
linux> make
```

WARNING: Do not let the Windows WinZip program open up your `.tar` file (many Web browsers are set to do this automatically). Instead, save the file to your Linux directory and use the Linux `tar` program to extract the files. In general, for this class you should NEVER use any platform other than Linux to modify your files. Doing so can cause loss of data (and important work!).

4 Description

4.1 Reference Trace Files

The `traces` subdirectory of the handout directory contains a collection of *reference trace files* that we will use to evaluate the correctness of the cache simulator you write in Part A. The trace files are generated by a Linux program called `valgrind`. For example, typing

```
linux> valgrind --log-fd=1 --tool=lackey -v --trace-mem=yes ls -l
```

on the command line runs the executable program “`ls -l`”, captures a trace of each of its memory accesses in the order they occur, and prints them on `stdout`.

Valgrind memory traces have the following form:

```
I 0400d7d4,8
M 0421c7f0,4
L 04f6b868,8
S 7ff0005c8,8
```

Each line denotes one or two memory accesses. The format of each line is

```
[space]operation address,size
```

The *operation* field denotes the type of memory access: “I” denotes an instruction load, “L” a data load, “S” a data store, and “M” a data modify (i.e., a data load followed by a data store). There is never a space before each “I”. There is always a space before each “M”, “L”, and “S”. The *address* field specifies a 64-bit hexadecimal memory address. The *size* field specifies the number of bytes accessed by the operation.

4.2 Writing a Cache Simulator

For this lab, you will write a cache simulator in `csim.c` that takes a valgrind memory trace as input, simulates the hit/miss behavior of a cache memory on this trace, and outputs the total number of hits, misses, and evictions.

We have provided you with the binary executable of a *reference cache simulator*, called `csim-ref`, that simulates the behavior of a cache with arbitrary size and associativity on a valgrind trace file. It uses the LRU (least-recently used) replacement policy when choosing which cache line to evict.

The reference simulator takes the following command-line arguments:

```
Usage: ./csim-ref [-hv] -s <s> -E <E> -b <b> -t <tracefile>
```

- `-h`: Optional help flag that prints usage info
- `-v`: Optional verbose flag that displays trace info
- `-s <s>`: Number of set index bits ($S = 2^s$ is the number of sets)
- `-E <E>`: Associativity (number of lines per set)
- `-b `: Number of block bits ($B = 2^b$ is the block size)
- `-t <tracefile>`: Name of the valgrind trace to replay

The command-line arguments are based on the notation (s , E , and b) from page 425 of the CS:APP3e (Chinese ver.) textbook. For example:

```
linux> ./csim-ref -s 4 -E 1 -b 4 -t traces/yi.trace
hits:4 misses:5 evictions:3
```

The same example in verbose mode:

```
linux> ./csim-ref -v -s 4 -E 1 -b 4 -t traces/yi.trace
L 10,1 miss
M 20,1 miss hit
L 22,1 hit
S 18,1 hit
L 110,1 miss eviction
L 210,1 miss eviction
M 12,1 miss eviction hit
hits:4 misses:5 evictions:3
```

Your job is to fill in the `csim.c` file so that it takes the same command line arguments and produces the identical output as the reference simulator. Notice that this file is almost completely empty. You'll need to write it from scratch.

4.3 Programming Rules

- Include your name and ECNU student ID in the header comment for `csim.c`.
- Your `csim.c` file must compile without warnings in order to receive credit.
- Your simulator must work correctly for arbitrary s , E , and b . This means that you will need to allocate storage for your simulator's data structures using the `malloc` function. Type "man malloc" for information about this function.
- For this lab, we are interested only in data cache performance, so your simulator should ignore all instruction cache accesses (lines starting with "I"). Recall that `valgrind` always puts "I" in the first column (with no preceding space), and "M", "L", and "S" in the second column (with a preceding space). This may help you parse the trace.
- To receive credit, you must call the function `printSummary`, with the total number of hits, misses, and evictions, at the end of your `main` function:

```
printSummary(hit_count, miss_count, eviction_count);
```

- For this lab, you should assume that memory accesses are aligned properly, such that a single memory access never crosses block boundaries. By making this assumption, you can ignore the request sizes in the `valgrind` traces.

5 Working on the Lab

We have provided you with an autograding program, called `test-csim`, that tests the correctness of your cache simulator on the reference traces. Be sure to compile your simulator before running the test:

```
linux> make
linux> ./test-csim
```

Points	(s,E,b)	Your simulator			Reference simulator			
		Hits	Misses	Evicts	Hits	Misses	Evicts	
6	(1,1,1)	9	8	6	9	8	6	traces/yi2.trace
6	(4,2,4)	4	5	2	4	5	2	traces/yi.trace
6	(2,1,4)	2	3	1	2	3	1	traces/dave.trace
6	(2,1,3)	167	71	67	167	71	67	traces/trans.trace
6	(2,2,3)	201	37	29	201	37	29	traces/trans.trace
6	(2,4,3)	212	26	10	212	26	10	traces/trans.trace
6	(5,1,5)	231	7	0	231	7	0	traces/trans.trace
8	(5,1,5)	265189	21775	21743	265189	21775	21743	traces/long.trace

50

TEST_CSIM_RESULTS=50

For each test, it shows the number of points you earned, the cache parameters, the input trace file, and a comparison of the results from your simulator and the reference simulator.

Here are some hints and suggestions for working on Part A:

- Do your initial debugging on the small traces, such as `traces/dave.trace`.
- The reference simulator takes an optional `-v` argument that enables verbose output, displaying the hits, misses, and evictions that occur as a result of each memory access. **You are NOT required to implement this feature in your `csim.c` code, but we strongly recommend that you do so.** It will help you debug by allowing you to directly compare the behavior of your simulator with the reference simulator on the reference trace files.
- We recommend that you use the `getopt` function to parse your command line arguments. You'll need the following header files:

```
#include <getopt.h>
#include <stdlib.h>
#include <unistd.h>
```

See “`getopt(3)` — Linux manual page” for details.

- Each data load (L) or store (S) operation can cause at most one cache miss. The data modify operation (M) is treated as a load followed by a store to the same address. Thus, an M operation can result in two cache hits, or a miss and a hit plus a possible eviction.

6 Evaluation

This section describes how your work will be evaluated. The full score is 100 points:

- Correctness: 50 points
- Coding style: 20 points
- Report: 30 points

6.1 Correctness (50 pts)

For correctness, we will run your cache simulator using different cache parameters and traces. There are eight test cases, each worth 6 points, except for the last case, which is worth 8 points:

```
linux> ./csim -s 1 -E 1 -b 1 -t traces/yi2.trace
linux> ./csim -s 4 -E 2 -b 4 -t traces/yi.trace
linux> ./csim -s 2 -E 1 -b 4 -t traces/dave.trace
linux> ./csim -s 2 -E 1 -b 3 -t traces/trans.trace
linux> ./csim -s 2 -E 2 -b 3 -t traces/trans.trace
linux> ./csim -s 2 -E 4 -b 3 -t traces/trans.trace
linux> ./csim -s 5 -E 1 -b 5 -t traces/trans.trace
linux> ./csim -s 5 -E 1 -b 5 -t traces/long.trace
```

You can use the reference simulator `csim-ref` to obtain the correct answer for each of these test cases. During debugging, use the `-v` option for a detailed record of each hit and miss.

For each test case, outputting the correct number of cache hits, misses and evictions will give you full credit for that test case. Each of your reported number of hits, misses and evictions is worth 1/3 of the credit for that test case. That is, if a particular test case is worth 6 points, and your simulator outputs the correct number of hits and misses, but reports the wrong number of evictions, then you will earn 4 points.

6.2 Coding Style (20 pts)

There are 20 points for coding style. These will be assigned manually by the TAs. Style guidelines can be found [here](#).

6.3 Report (30 pts)

You should use Markdown to write your report in the file `README.md`. Your report should contain the following sections:

- Title, your name and student ID.
- Implementation of the cache simulator, including:

- Data structure.
- Program flow.
- Function description and calling relationship.

It is recommended to use pseudo-code, flowcharts, etc. for description.

- The output of the autograding program `test-csim`.
- Summary: what you have learned from this lab.

7 Handing in Your Work

When you finish your assignment, please make sure your work (code and report) is pushed to `homework04` branch of your repository in Shuishan before the deadline.

```
linux> git push <REMOTE_NAME> master:homework04
```

IMPORTANT: Do not create the handin tarball on a Windows or Mac machine, and do not handin files in any other archive format, such as `.zip`, `.gzip`, or `.tgz` files.